Evolving Character Relationships Using HMM and LSTM

Sushmita Gopalan　　　　　**Haylee Ham**　　　　　**Chelsea Ernhofer**

Abstract

The purpose of this project is to use small book synopses to predict cooperative or uncooperative character relationships. This is based off of work by Snigdha Chaturvedi et al. (2015) in which evolving character relationships are modeled by training a second order Hidden Markov Model on key text features. We attempt to capture similar features as Chaturvedi et al. through feature extraction, however, we use different libraries to do so. Our group recreates the modeling and extends it by also using LSTM Neural Networks to capture character relationships. Results show a discrepency between our work and previous work, most likely due to different in feature extraction and implementation of first order HMM rather than a second order model. Our LSTM model performs well, however, certain limitations in our construction of the LSTM lead us to believe that we overstate model metrics.

1 Introduction

Computational narrative studies is a field within natural language processing which seeks to understand, summarize, and generate stories. An increased ability to form narratives is not only useful when applied to literature or film, but can increase a machine's ability to communicate effectively with human users. Identification of character relationships is an important method used in computational narrative studies to identify and summarize individual character roles within a narrative and thus more fully understand the literature in question.

Current research has developed multiple methods in attempts to effectively understand stories through analysis of the characters within them. Both supervised and unsupervised methods have been applied to tasks such as role and relationship extraction and social network modeling. Our work seeks to identify evolving character relationships through supervised methods.

2 Related Work

The body of this work is based off of a paper by Snigdha Chaturvedi et al. (2015) in which researchers attempt to model the specific and dynamic relationships between characters. This goal is divergent from other character relationship projects which seek to simply predict one static relationship between characters or identify basic character roles within a narrative. This is useful since the relationship between two given characters is likely to change throughout the progression of the novel. Chaturvedi et al. used novel summaries from the website SparkNotes as data. From these summaries, sentences were extracted which contained information on a pair of characters. Human readers went through these sentences and determined whether the relationship between the two main characters mentioned was cooperative or noncooperative. This became the target variable. For the predictors, text was then preprocessed and features such as part of speech tags and dependency parses were extracted. From there, Chaturvedi et al. engineered two types of features: content based features and transition features. Content based features contain information concerning the verbs and adverbs that two characters complete together, separately, or one to another. Content features also include semantic parses which employs a number of frames to the input sentences. Transition features indicate whether the relationship between characters changes between the different extracted sentences.

Chaturvedi et al. use logistic regression and decision trees as an unstructured baseline. In these models, each sentence was fed through independently of the sequence to which it belongs and the relationship between the characters in that one sentence was predicted as either cooperative or noncooperative. A second-order Markov Model was then used to capture the potentially changing relationships between characters over time. Precision, recall, and F1 score were used as evaluation metrics. Results showed that the second order Markov Model had the highest F1 score out of all methods used with a measure of 60.76 compared to 48.54 generated by the decision tree and 51.48 from the logistic regression.

3 Data

Data from the Chaturvedi et al. paper was used for this project. Data consisted of 100 sequences comprising almost 800 single sentences. Half of these sequences were fully annotated, that is, they included indications of which characters were central in each sequence and marked their reference in each individual sentence and all sentences in each sequence had a label. For the other 50 sequences which did not include such annotation, annotations were added manually after reading through the sentences and external plot descriptions.

4 Methods

We attempt to recreate the majority of the content based features found in the Chaturvedi et al. paper, namely verb based interactions between characters and counts of positive and negative spanning words between characters. We then also implement a decision tree and logistic regression as baseline indicators. A first order Hidden Markov Model is used to generate predictions based off of complete sequences rather than single sentences. Finally, we extend the work of Chaturvedi et al. by including a Neural Network model to also predict character relationships, allowing for evolution of said relationships.

4.1 Feature Extraction

Positive and negative verb interaction features were extracted for three separate instances: characters verb together, one character verbs another, and characters verb separately. In order to engineer these features, part of speech tagging and dependency parsing were run to capture the struc-

ture of the sentence and individual word roles. For characters verb together, verbs were included if they were ancestors of both characters in a sentence. For one character verbs another, a verb was included if it was between an object and subject and the object and subject were the two main characters of that sentence. Verbs that were an ancestor to either of the characters and weren't already included in another set (either characters verb together or one character verbs another) were added to the single character verbs set. Sentiment analysis was then done on all of the verbs included and a binary indicator was created to identify whether characters negatively or positively verbed. In order to incorporate frame semantic parsing, our group used the frames features created by Chaturvedi et al. Chaturvedi et al. manually compiled frames that were positive or negative in nature and then included binary features indicating whether or not positive, neutral, or negative frames fired between the two characters in question. Two additional features were included that indicate the count of positive and negative words that appear in the span of words between mentions of the two characters.

4.2 Baseline Modeling

In adherence to the Chaturvedi et al. paper, we use a logistic regression and decision trees as baseline models. Both these methods are used for classification and allow for varied levels of nonlinearity.

4.3 HMM

As hypothesized in Chaturvedi et al., relationship sequence prediction is a structured problem. The baseline methods of logistic regression and decision trees treat it as a flat classification problem. The Hidden Markov Model (HMM) is the method that Chaturvedi et al. proposes be used to address the structured nature of this problem. An HMM is a generative model, where a sequence of observable events is generated by a sequence of latent or hidden states. The model we use is a first-order HMM, i.e. we assume that transitions between latent states have the form of a first order Markov chain. The Viterbi algorithm is used for decoding probabilities.

4.4 LSTM

We further propose fitting a long short-term memory network. Similarly to HMMs, LSTM networks can make sequence predictions taking in

the entire sequence of inputs and outputting a sequence of predictions about the sequences. The architecture of the LSTM network is made up of two hidden LSTM layers as well as an output layer, activated by the sigmoid function. Binary cross entropy loss was used to train the network. In order to avoid overfitting, both early stopping and dropout were incorporated into the model. There are two dropout layers in the architecture, one in between the LSTM layers and one in between the second LSTM layer and the output layer. In training the LSTM network, several hyperparameters were tuned: the number of epochs, the batch size, dropout rate, the number of hidden layers, and and the number of dropout layers. The following values were tested as hyperparameters: 1 and 2 hidden LSTM layers, 1 and 2 dropout layers, 10% and 20% dropout rates, between 4 and 20 epochs, and between 5 and 20 batch size. At the conclusion of training, 2 LSTM hidden layers, 2 dropout layers, 20% dropout rate, 10 epochs, and a batch size of 10 were chosen as the hyperparameters yielding the most accurate results.

5   Experimental Setup

All work and data can be found in our group's github page https://github.com/sushmitavgopalan16/Very-Fun-Characters

5.1   Data

Data can be found on Snigdha Chaturvedi's professional website at https://sites.google.com/site/snigdhac/academics. Our team downloaded the data from the 2015 paper 'Modeling Evolving Relationships between Characters in Literary Novels'. Data is formatted in such a way that there exists a separate csv file for each sequence. We combined these files into one, including character IDs for each sequence in the final data frame. For the partially annotated data, we included our manual annotations at this point in the process. Sentences that had no label (were neither cooperative nor noncooperative) were discarded, leaving only two possible classes: cooperative and noncooperative. This left us with 603 sentences and 100 sequences.

5.2   Feature Extraction

The SpaCy library was used to perform part of speech tagging and dependency parsing. We then wrote an algorithm to identify ancestors and de-

termine verbs included in an object-verb-subject clause. In order to create the positive/negative span features, the vaderSentiment package was used. We created an algorithm that searched for the two characters in question and then found the sentiment of all words that occur in between mentions of the two characters. The total number of positive and negative words found in the spanned words in the sentence where then summed and the two features were created for each sentence.

5.3   Modeling

The Decision Tree and Logistic Regression were performed using the sklearn library. Data were partitioned into a training and test set (with ratio 80:20) using the train test split functionality included in the sklearn library.

For the HMM, we used an implementation found https://github.com/larsmans/seqlearn as a guideline and corrected a minor error in the calculation of emission probabilities. Python's scikitlearn package no longer supports an API for a hidden markov model.

We split the data into training and validation sets in an 80-20 ratio. We split the sequences, rather than the sentences themselves, in order to ensure that no sequences get split up over the training and testing subsets. We ran a 100-fold cross-validation to reduce overfitting on our relatively small set of sequences.

In training the LSTM network, the data was transformed so that rather than feeding in one sentence at a time, each observation was a concatenation of all of the features for every sentence in the sequence. Similarly, the target variable was a sequence of indicators concerning whether each sentence in the sequence portrayed the two characters in question cooperating or not. Padding was done using the pad sequences functionality of the Keras package.The data was then fed into a network created in the Keras deep learning framework. Data were split into a training and test set using the sklearn library's train test split functionality. The split was 80% training and 20% test set.

5.4   Model Evaluation

For each model used, three statistics were calculated: precision, recall, and F1 score. Overall model success was based on F1 score since it is a summary of both precision and recall and was the primary evaluation metric used in the Chaturvedi et al. paper.

6 Results and Analysis

Model	Precision	Recall	F1 Score
LR	40.42	49.54	44.55
DT	55.87	55.41	55.69
HMM	39.16	49.95	43.88
LSTM	83.34	81.75	82.49

Table 1. Evaluation of Classification Models- Current work

Table 1 contains the results of each of the four models used in analysis. The models which perform the worst are the logistic regression (F1 of 0.445) and the first order markov model (F1 of 0.438). Examining model metrics per class for the first order HMM, we see that none of the sentences labeled as describing non-cooperative relationships were identified by our model. This can likely be attributed to the class imbalance in our data, where only about one in five sentences are labeled as non-cooperative. The decision tree model performs slightly better with an F1 score of 0.556. The LSTM performs the best out of these models, however, padding may have influence these results. This potentially biased result is one limitation of our work and likely caused the model evaluation metrics to be much higher than they are in reality.

Model	Precision	Recall	F1 Score
LR	54.98	51.83	50.05
DT	63.21	55.79	55.77
Order 1 HMM	62.45	63.40	62.84
Order 2 HMM	66.49	65.97	66.22

Table 2. Evaluation of Classification Models- Chaturvedi et al. 2015

Table 2 contains the results from the Chaturvedi et al. 2015 paper. Results for almost all models differ greatly between our work and the work done by Chaturvedi et al. This is likely due to differences in the feature extraction methodology. The exact methods used by Chaturvedi et al. were not disclosed in the 2015 paper and the specific library used for part of speech tagging, depedency parsing, and framing has been officially deprecated. One interesting note is that the decision tree models have very similar scores. This is likely due to the fact that cross validation was not used in the creation of the decision tree.

In comparing our results to those of Chaturvedi et al., it must be taken into account that Chaturvedi et al. included a semi-supervised framework.

Specifically, the study included some data that was not labeled as cooperative or non cooperative and Chaturvedi et al. trained a segmentation model to make predictions about the label of the unsupervised sentences. These labels were not provided by Chaturvedi et al. As previously stated, our study incorporates a fully supervised framework by removing these unlabeled observations. In total, 189 observations did not have labels and were removed.

7 Limitations and Conclusion

In attempting to recreate the features described in the Chaturvedi et al. study, we found that there were several aspects of the feature creation process that were not clearly laid out in the paper. In light of this and the fact that one of the packages mentioned in the Chaturvedi et al. paper had been deprecated, our group had to make several choices in how to proceed. Different packages and software were used in both creating the features and estimating the models. As mentioned above, in creating the LSTM network, padding was added to both the input and output sequences in order to maintain a consistent shape of the data. This padding altered the ratio of cooperative and noncooperative labels, since in the LSTM the cooperative labels were coded as 1 and the noncooperative as 0. This may account for some of the high accuracy in the LSTM as the other models we estimated suffered from class imbalance issues.

Another primary challenge we faced was the paucity of training data. No naturally occurring data exists that labels the nature of relationships between characters in single sentences. Crowdsourcing annotations through a platform like Amazon's MTurk could be a next step in collecting more data that could improve our models.

For the decision tree and logistic models, further cross-validation could be valuable in helping the model generalize better. Experimenting with second and third order Markov models, thus incorporating more of the history of the relationship of the pair of characters could potentially improve this model significantly.

8 References

Chaturvedi, Snigdha, et al. "Modeling dynamic relationships between characters in literary novels." arXiv preprint arXiv:1511.09376 (2015).